


Napredni modeli i baze podataka

Predavanja

9. Spatio-temporal Databases

Studenii 2008.

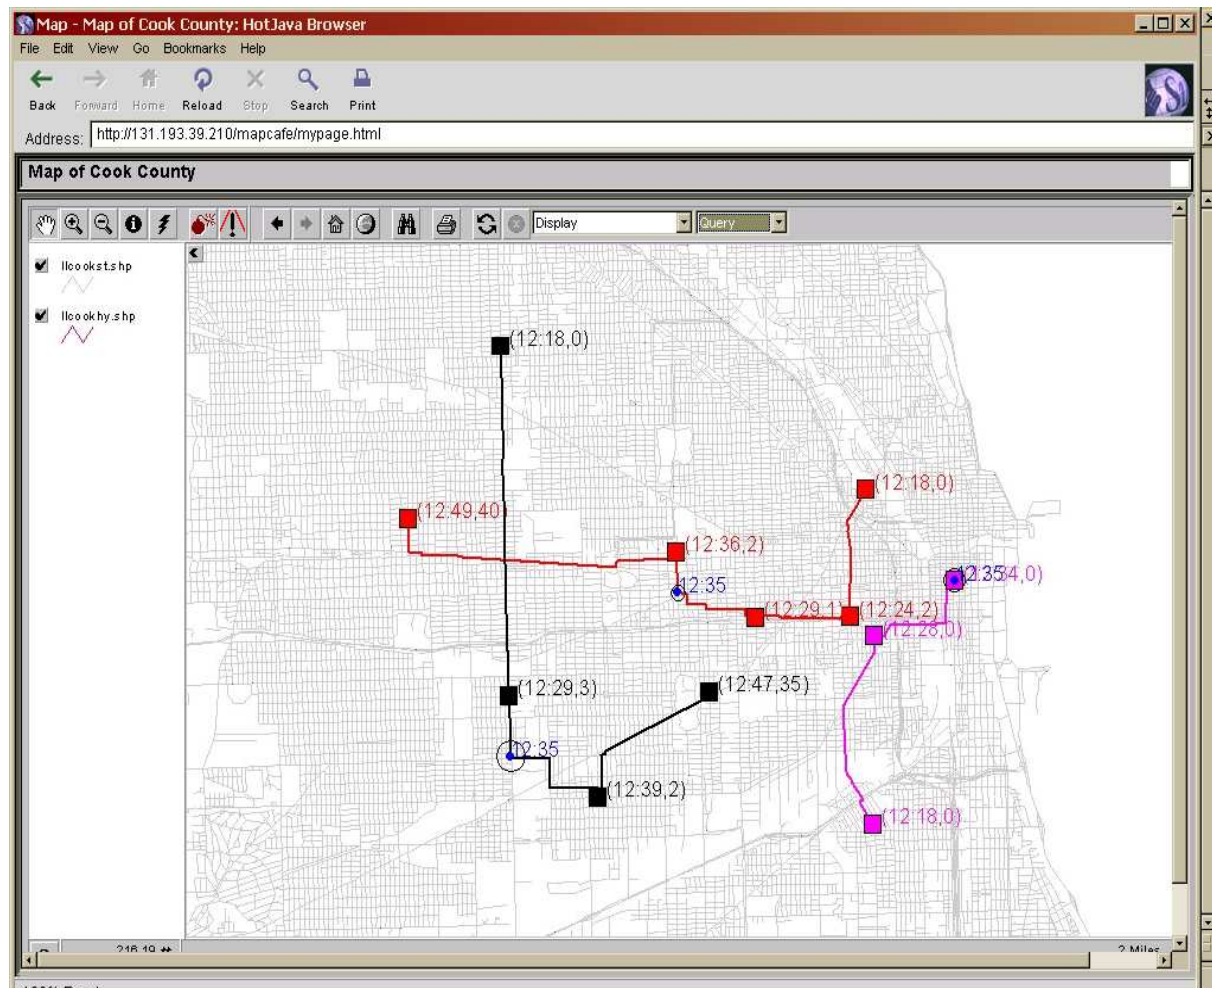


Pregled

- Motivation Examples
- Definitions
- Managing Time in Standard DB
- Time Domain and Time Dimensions
- Taxonomy of Temporal Databases
- Temporal Structured Query Language - TSQL2
- Moving Object Databases
- Examples

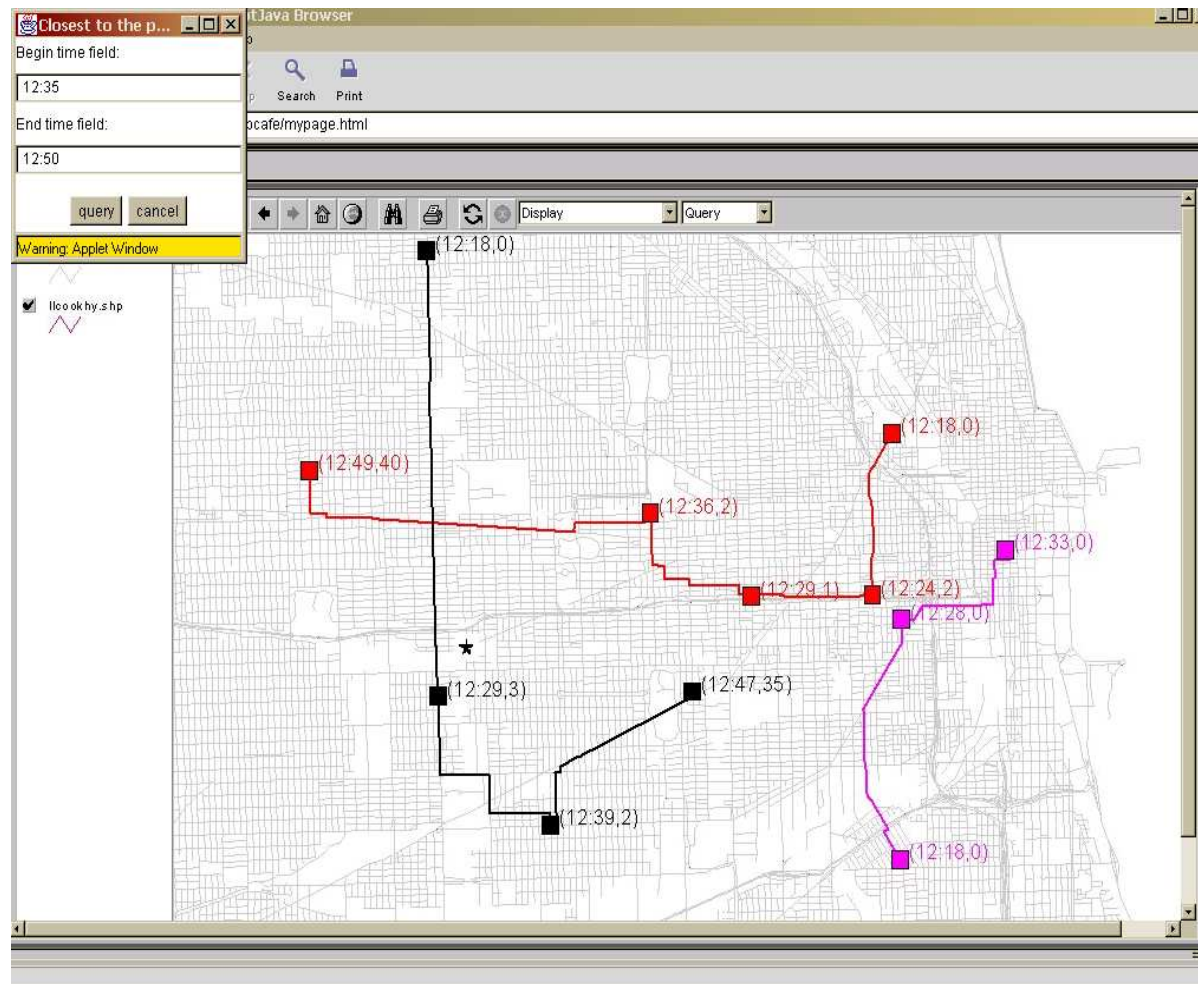
Examples

Q_1 : Show all vehicle's locations at 12:35



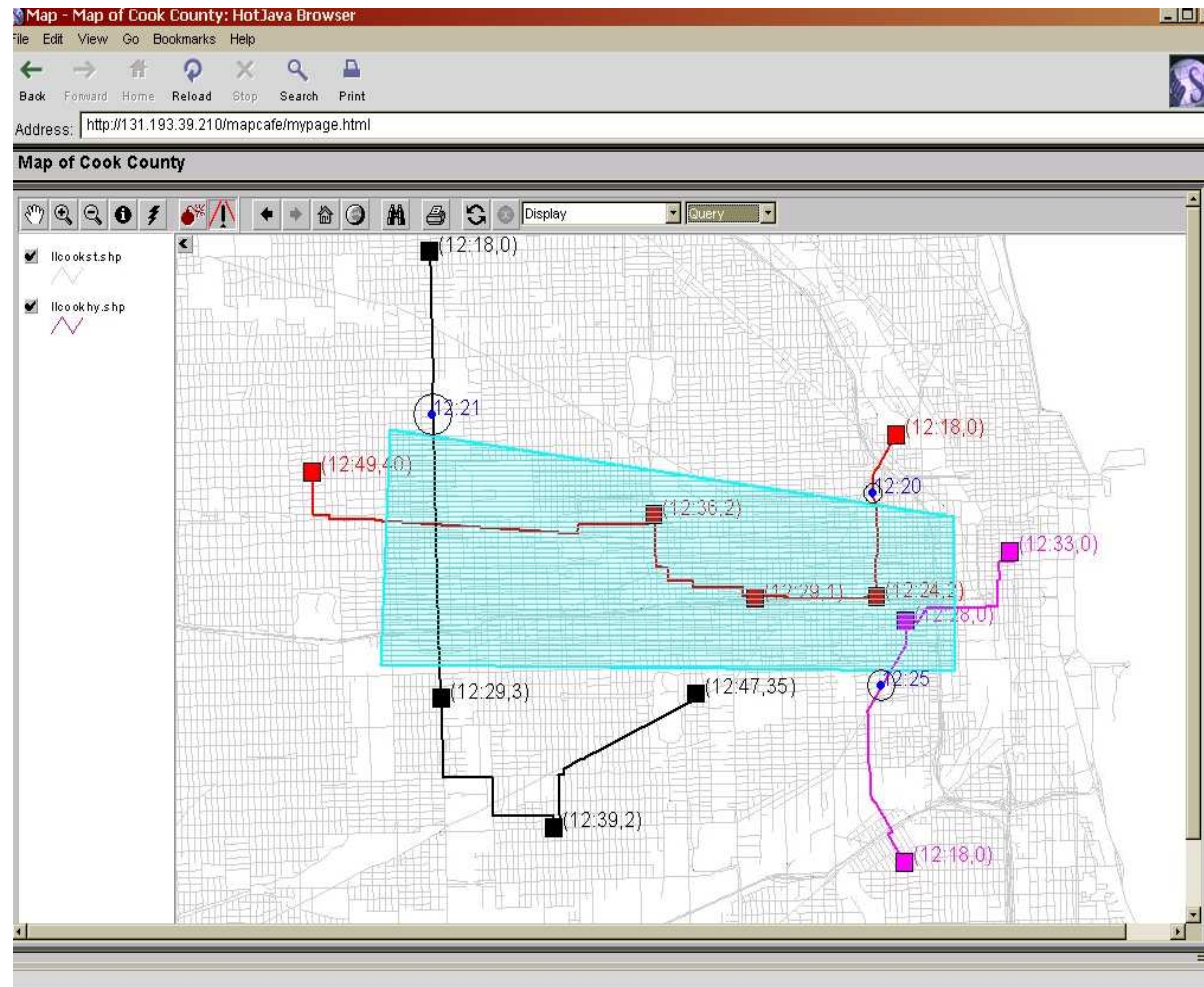
Examples

Q_3 : Which vehicle will be closest to the “star” between 12:35 and 12:50?



Examples

Q_3 : *When will each vehicle enter the specified sector?*



Definitions

- A **spatio-temporal** data base is (a new type of) database system that manages spatio-temporal objects and supports corresponding query functionalities
- A **spatio-temporal object** is a kind of object that dynamically changes spatial location and/or extents along with time.
 - A typical example of a spatiotemporal object is a **moving object** whose location continuously changes.

Jespersen and Fitz-Randolph, *From Sundials to Atomic Clocks*

Managing Time in Standard DB

- Databases managed by standard DBMS reflect current state of the world as far as known in the database
- UPDATE → previous state is lost. However, keeping track of history needed by many applications. In standard DBMS possible if **application** manages time itself.
- Standard DBMS/SQL offer limited temporal support in the form of data types
 - **date** (a particular day from a year in the range 1 through 9999 A.D.)
 - **time** (a particular second within a range of 24 hours)
 - **timestamp** (a particular fraction of a second of a particular day)
 - **interval** (a directed relative duration of time – time interval of known length with unspecified start and end instants)

Managing Time in Standard DB

- Example:

```
parcel (city: string, number: integer, geometry: polygon)
```

Add attributes for managing time:

```
parcel (city: string, number: integer, geometry: polygon,  
from: timestamp, to: timestamp)
```

- Disadvantages:

- Temporal semantic, operations and integrity constraints built into application(s)
- Difficult, error-prone, complex queries
- Inefficient query execution

How to Implement ST Applications ?

- **Use built-in SQL data types and build temporal support in application(s)**
 - Disadvantages
 - ... *discussed on the previous slide*
- **Implement an abstract data type (ADT) for time**
 - Disadvantages
 - Not possible using a pure relational model
 - ... *similar to the previous approach*
- **Extend non-temporal model to temporal data model**
 - Extend non-temporal schema with special temporal attributes
 - Extend algebra and query language with additional operators to express a *temporal join, temporal selection, temporal projection, etc.*
 - Disadvantages
 - Proposed extended temporal usually concentrate on special features:
 - Temporal data structures
 - Query language design
 - Temporal algebra
 - Temporal integrity constraints

How to Implement ST Applications ?

- **Generalize non-temporal data model to a temporal data model**

All three components have to be generalized:

- Data structures, operations and integrity constraints
Type or schema of objects is not simply extended, but a new, ***simple*** and ***orthogonal*** concepts needs to be found
- ***simple***: easily implementable
- ***orthogonal***:
 - concept is ***not restricted*** to specific constructs of the data model (tuples, attributes, ...)
 - User should decide which granularity of data and even constructs of the model (types, collections, integrity constraints or even DB) shall be temporal
 - Temporal operations (including updates) shall refer to this new concept

The Goal of Spatio-Temporal Research

- Integrate spatio-temporal concepts deeply into DBMS data model and query language
- Extend the system implementation accordingly for efficient execution
- Built-in temporal support provides :
 - higher-fidelity data modeling
 - More efficient application development
 - Potential increase in performance

Time

- *It's presented everywhere, but occupies no space*
- *We can measure it, but we can't see it, touch it, get rid of it, or put it in a container*
- *Everyone knows what it is and uses it every day, but no one has been able to define it.*

Jespersen and Fitz-Randolph, *From Sundials to Atomic Clocks*

The Time Domain

- Time is generally perceived as a one-dimensional space extending from the past to the future:




- There are some options:
 - ***bounded/infinite***
 - ***discrete/dense/continuous***
 - Discrete models are isomorphic to the natural numbers or integers
 - Each natural number corresponds to an atomic time interval – ***chronon***
 - Chronons can be grouped into larger units – ***granules*** (hours, days, ...)
 - Dense models are isomorphic to either the rationals or the reals
 - between any two instants of time another instant of time exists
 - Continuous models are isomorphic the real numbers
 - Each real number corresponds to a ***point in time***
 - ***absolute/relative***
 - November 20, 2008, 15:45
 - two weeks

Temporal Data Types

- Previous concepts of time can be captured in a number of data types:
 - **instant**
A particular *chronon* on the time line in the discrete model or a point on the timeline in a continuous model
 - **period**
An anchored interval on the time line
 - Unlike standard SQL temporal types, period is **not** ordered – there is only a **partial order** between periods
 - Although is not part of SQL standard, period is relatively easy to simulate with `date/time`.
 - **periods**
A set of disjoint anchored intervals on the timeline - also called a *temporal element* in the literature.

Relationships/Predicates Between Two Periods

- There are 13 possible relationships/predicates between two periods:



Relationships/Predicates Between Two Periods

- Semantics of Allen's operators


Comparison Predicate	Equivalent Predicates on Endpoints
I_1 before I_2	$end(I_1) < begin(I_2)$
I_1 after I_2	$end(I_2) < begin(I_1)$
I_1 during I_2	$(begin(I_1) > begin(I_2) \wedge end(I_1) \leq end(I_2)) \vee$ $(begin(I_1) \geq begin(I_2) \wedge end(I_1) < end(I_2))$
I_1 contains I_2	$(begin(I_2) > begin(I_1) \wedge end(I_2) \leq end(I_1)) \vee$ $(begin(I_2) \geq begin(I_1) \wedge end(I_2) < end(I_1))$
I_1 overlaps I_2	$begin(I_1) < begin(I_2) \wedge end(I_1) > begin(I_2) \wedge end(I_1) < end(I_2)$
I_1 overlapped_by I_2	$begin(I_2) < begin(I_1) \wedge end(I_2) > begin(I_1) \wedge end(I_2) < end(I_1)$
I_1 meets I_2	$end(I_1) = begin(I_2)$
I_1 met_by I_2	$end(I_2) = begin(I_1)$
I_1 starts I_2	$begin(I_1) = begin(I_2) \wedge end(I_1) < end(I_2)$
I_1 started_by I_2	$begin(I_1) = begin(I_2) \wedge end(I_2) < end(I_1)$
I_1 finishes I_2	$begin(I_1) > begin(I_2) \wedge end(I_1) = end(I_2)$
I_1 finished_by I_2	$begin(I_2) > begin(I_1) \wedge end(I_1) = end(I_2)$
I_1 equals I_2	$begin(I_1) = begin(I_2) \wedge end(I_1) = end(I_2)$

Time Dimensions


- In the context of databases, two **orthogonal** time dimensions are of general interests:
 - **Valid time**
Time in real world when an event occurs or a fact is valid, independent of the recording of that event/fact in some database.
 - **Transaction time**
Time when a change is recorded in the database, or the time interval during which particular state of the database exists.

Time Dimensions


snapshot relation




valid-time relation



transaction time relation



bitemporal relation



Time Dimensions

- Valid time and transaction time are:

- ***Orthogonal***
- ***but, not homogeneous***

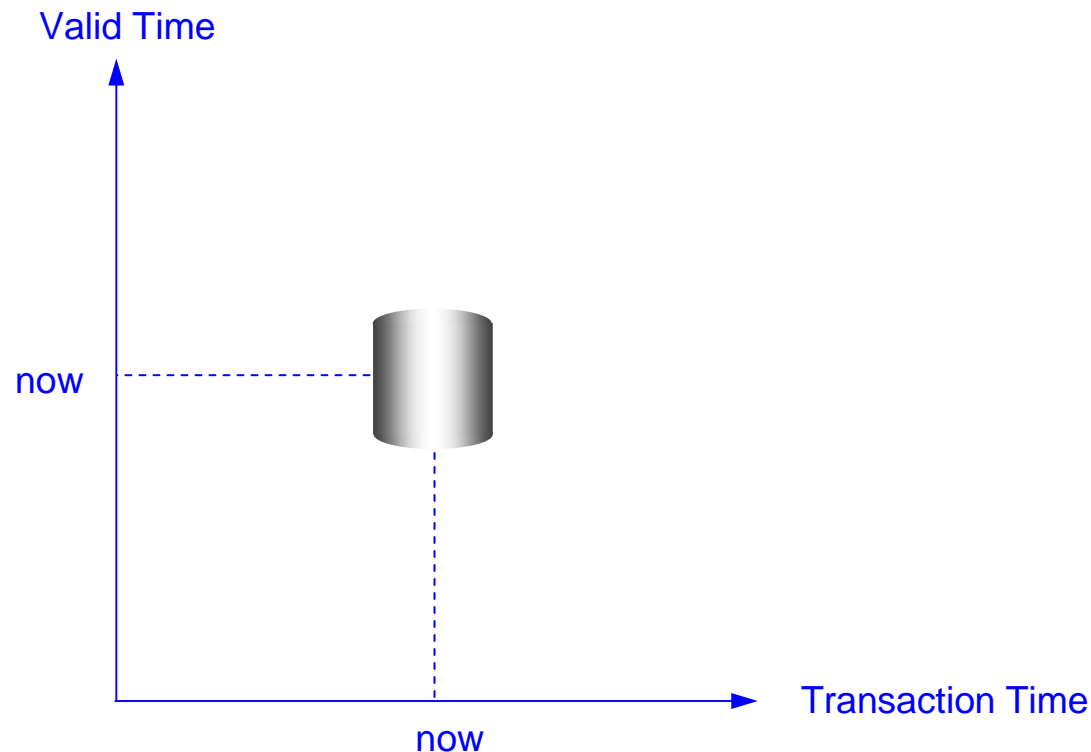
Transaction time has different semantic than valid time.

*Valid time may be extended into the future, transaction time is defined only until **now**.*

Taxonomy of Temporal Databases

■ *Snapshot Database*

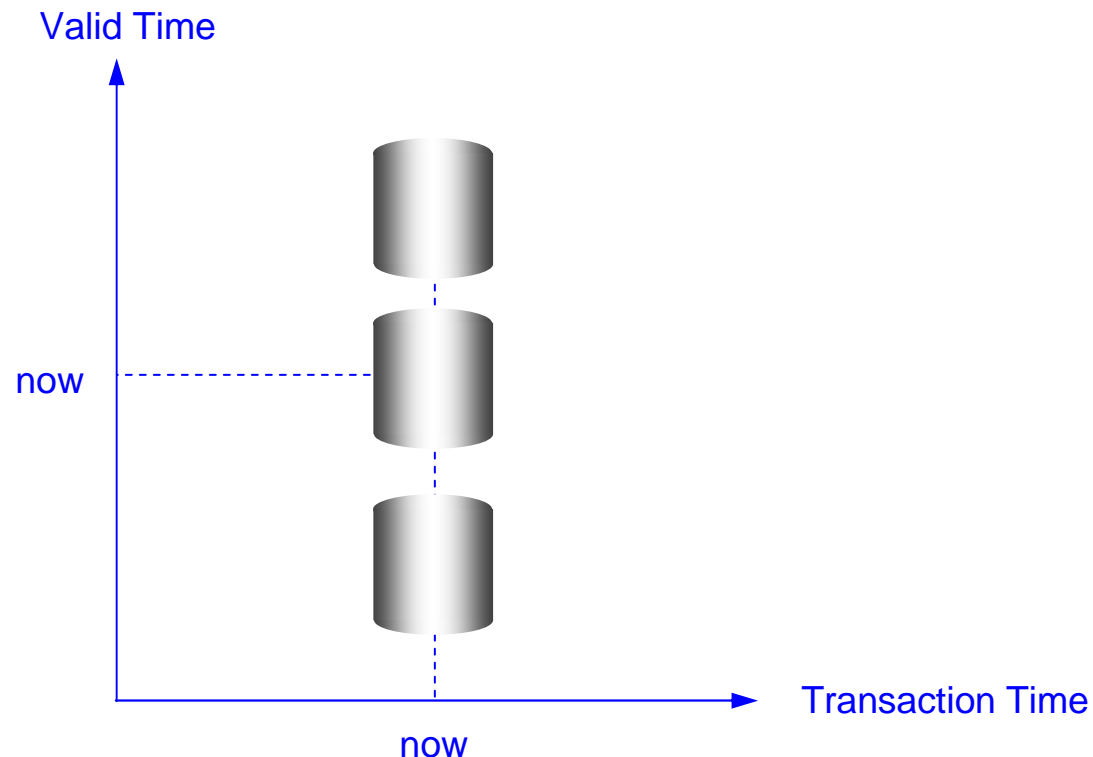
- Captures a single snapshot of the real world, usually current one
- This is what currently available commercial DBMS support
- Modifying the state of DB is done by INSERT, DELETE or UPDATE operations
- Past states of the database are overwritten



Taxonomy of Temporal Databases

■ *Historical Database*


- Records the history of data with respect to real world, i.e. records database state along the **valid time** axis.
- The user must supply the valid time of a fact
- Data is still physically deleted when corrected



Taxonomy of Temporal Databases

■ **Rollback Database**

- Records the changes to DB itself, i.e. records database state along the **transaction time** axis.
- An *append-only* DB, i.e., no data is ever physically deleted




We consider **linear transaction time** only.

An alternative - **branching transaction time**, provides a useful model for **versioning (long transaction)**.

Taxonomy of Temporal Databases

■ *Bitemporal Database*

- A combination of a historical and a rollback DB.
- Records database states with respect to both ***valid time*** and ***transaction time***



The Temporal Structured Query Language - TSQL2

- Based on the **Bitemporal Conceptual Data Model** (BCDM)

- On the 6th of January, the administration was informed that Lisa had started to work in the toys department on the 1st and was going to work there until the 15th.
- On the 10th it became known and entered into the database that Lisa had moved to the

<i>Name</i>	<i>Dept.</i>	<i>Time</i>
Lisa	Toys	{(6, 1), ..., (6, 15), ..., (9, 1), ..., (9, 15), (10, 1), ..., (10, 7), ..., (19, 1), ..., (19, 7), (<i>uc</i> , 1), ..., (<i>uc</i> , 7), (12, 14), ..., (12, 20), ..., (19, 14), ..., (19, 20), (<i>uc</i> , 14), ..., (<i>uc</i> , 16)}
Lisa	Books	{(10, 8), ..., (10, 13), ..., (19, 8), ..., (19, 13), (<i>uc</i> , 8), ..., (<i>uc</i> , 13), (10, 14), (10, 15), (11, 14), (11, 15)}
John	Books	{(12, 11), (12, 12), ..., (12, ∞), (13, 11), ..., (13, ∞), ..., (19, 11), ..., (19, ∞), (<i>uc</i> , 11), ..., (<i>uc</i> , ∞)}

Bitemporal space

- On the 12th it was decided that Lisa would move back to toys on the 14th and would stay there a while longer, until the 20th. Also a new employee John had started the day before in the books department.
- On the 20th, it was entered that Lisa had actually quit the company on the 16th.

The Temporal Structured Query Language - TSQL2

- An example of bitemporal relation:

```
CREATE TABLE parcel (  
    city:      STRING,  
    number:    INTEGER,  
    geometry:  POLYGON)  
AS VALID STATE DAY AND TRANSACTION
```

- **STATE*** – relation/DBMS records facts that are true over certain periods of time
- **DAY** – valid time granularity (in this example - one day)
- Transaction time granularity is system dependent (like milliseconds)

** An alternative clause **EVENT** records events that occurred at certain instants of time.*

The Temporal Structured Query Language - TSQL2

- There are six different kinds of relations in TSQL2:
 - Snapshot relations (no temporal support)
 - Valid-time state relations (... AS VALID [STATE])
 - Valid-time event relations (... AS VALID EVENT)
 - Transaction-time relations (... AS TRANSACTION)
 - Bitemporal state relations (... AS VALID STATE AND TRANSACTION)
 - Bitemporal event relations (... AS VALID EVENT AND TRANSACTION)

The Temporal Structured Query Language - TSQL2

- Display geometry of parcels (now or in the past):

```
SELECT SNAPSHOT p.geometry  
FROM parcel AS p;
```

- Display juridical geometry changes of parcel 456 in the City of Vienna:

```
SELECT geometry  
FROM parcel  
WHERE city = "Vienna"  
AND    number = 456;
```

The Temporal Structured Query Language - TSQL2

- Display juridical geometry changes of parcel 456 in the City of Vienna, during 2000:

```
SELECT p.geometry
VALID INTERSECT(VALID(p), PERIOD '[2000]' DAY)
FROM parcel AS p
WHERE p.city = "Vienna"
AND    p.number = 456;
```

- Display adjacent parcels of parcel 456 in the City of Vienna as of March 17, 1977:

```
SELECT ap.geometry
FROM parcel AS p, parcel AS ap
WHERE VALID(p) OVERLAPS DATE '[1977-03-17]'
AND    p.city = "Vienna"
AND    p.number = 456
AND    p.geometry.touches(ap.geometry);
```



Oracle Workspace Manager

- Wrapper/Simulator on top of Oracle DBMS
- **Black box** which cannot be changed (wrapped package)
- Does't support SQL-like temporal queries, but


```
EXECUTE dbms_wm.GoToDate( '06-APR-2002' );  
SELECT number, geometry  
FROM parcel_hist  
WHERE number = '1490/1';
```


The Temporal Structured Query Language - TSQL2

- Supports **discrete** spatio-temporal changes



- ... but not **continuous** spatio-temporal changes





Moving Objects Databases – Current Movement

- Moving objects spatio-temporal model - MOST
 - ***Trajectory location management prospective***, i.e. management the locations of a set mobile objects that are moving ***right now***
 - Supports current and expected future movement
 - FTL query language, based on ***future temporal logic***
 - Basically SQL extended with temporal operators
 - **UNTIL, NEXTTIME, EVENTUALLY, ALWAYS**
 - Restricted on moving points
 - Does not address any complex geometries
 - Query example:

```
SELECT id
FROM helicopters h
WHERE EVENTUALLY_WITHIN (inside (h, Valley), 10)
AND ALWAYS (inside (h, Valley), 2);
```

Moving Objects Databases – History of Movement

- Spatio-temporal data perspective
 - Geospatial objects stored in geospatial DB which change continuously over time
 - Spatio-temporal data types
 - Time dependent geometries with operations
 - ***moving point*** (people, cars, animals, planes, ships, ...)
 - ***moving region*** (forests, forest fires, countries, hurricanes, armies, ...)



Moving Objects Databases – History of Movement

- Assume some available operators

<i>moving(point)</i>	→	<i>line</i>	<i>trajectory</i>
<i>moving(region)</i>	→	<i>region</i>	<i>traversed</i>
<i>moving(α)</i>	→	<i>periods</i>	<i>deftime</i>
<i>moving(point) x moving(region)</i>	→	<i>moving(point)</i>	<i>intersection</i>
<i>moving(α) x instant</i>	→	<i>intime(α)</i>	<i>atinstant</i>
<i>intime(α)</i>	→	<i>instant</i>	<i>inst</i>
<i>intime(α)</i>	→	<i>α</i>	<i>val</i>
<i>periods(α)</i>	→	<i>int</i>	<i>duration</i>
<i>intime(α)</i>	→	<i>α</i>	<i>val</i>

moving - a type constructor that transforms a type α into a time dependent version of that type, *moving(α)*

Moving Objects Databases – History of Movement

- Examples

```
flight (id:STRING,from:STRING,to:STRING,route:mPoint)
weather (id : STRING, kind : STRING, area : mRegion)
```

All flights from Vienna that are longer than 5000 miles

```
SELECT id
FROM flight
WHERE from = "ZAG"
AND length(trajjectory(route)) > 5000;
```

Moving Objects Databases – History of Movement

- Examples

What was the route taken by flight 'OS 7052' ?

```
SELECT trajectory (route)
FROM flight
WHERE id = "OS 7052";
```

Which flights went through a snow storm ?

```
SELECT f.id
FROM flight f, weather w
WHERE w.kind = "snow storm"
AND duration(visits(f.route, w.area)) > 0
```

Which flights were in snow storm for more than 5 minutes ?

```
SELECT f.id
FROM flight f, weather w
WHERE w.kind = "snow storm"
AND duration(deftime(intersection(f.route, w.area))) > 300
```

Moving Objects Databases – History of Movement

- Implementation concept
ADT extension package to an extensible (OR, OO) DBMS
- Work program
 - Design a system of types and operations (*many sorted algebra*)
 - Which level of abstraction ?
 - Abstract model (infinite representation)
 - A continuous function from time into the Euclidean plane $f:\mathbb{R} \rightarrow \mathbb{R}^2$
 - Mathematically simple, elegant, and uniform, but not directly implementable
 - Discrete model (finite representation)
 - A polyline in the 3D space representing such a function
 - More complex and heterogeneous, but can be implemented
 - Design data structures for the types and algorithms for the operations
 - Implement DBMS extension package


Moving Objects Databases – History of Movement

- Discrete model (part)

region




abstract




discrete

a finite set of polygons,
each with polygonal
holes



edge-
disjoint

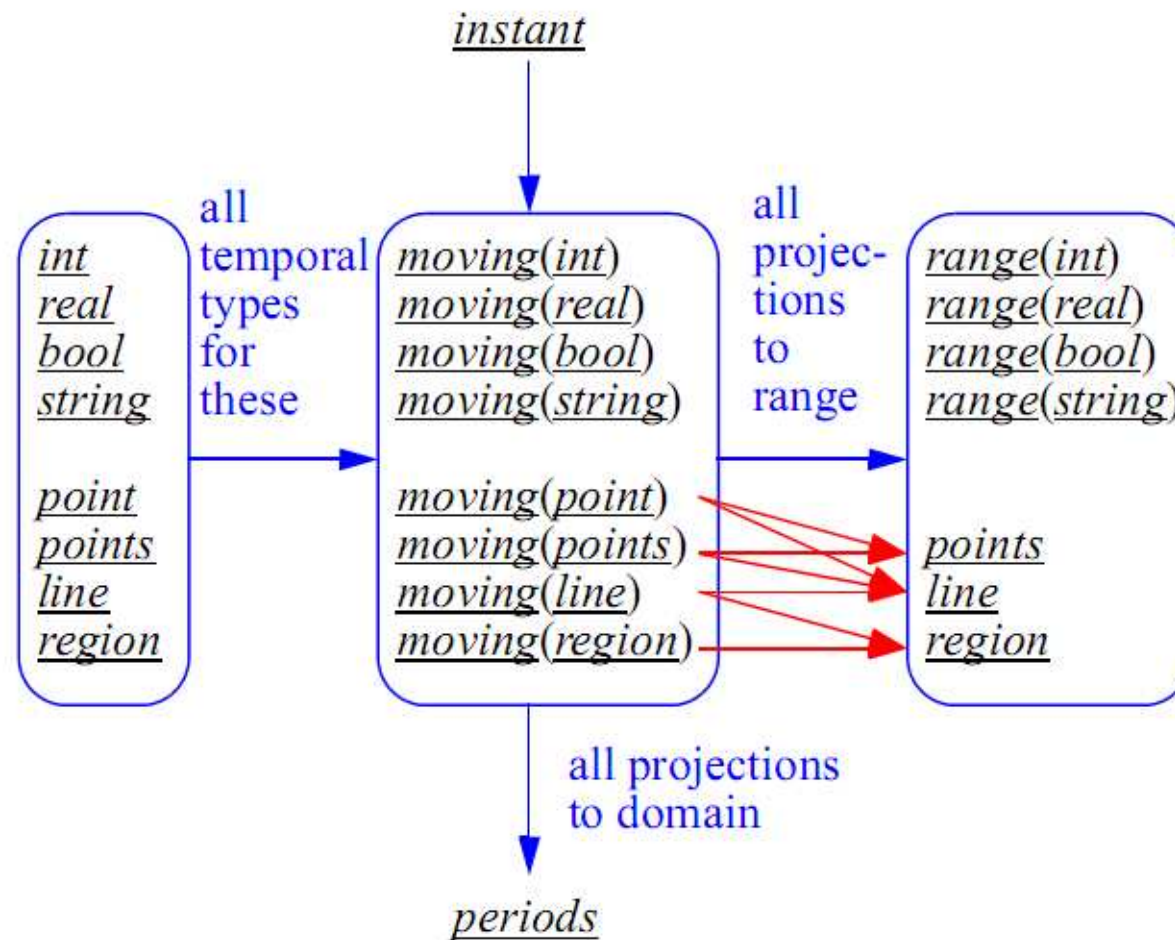


not edge-disjoint

$Seg = \dots$ (def. of line segments)
 $Cycle = \{S \subset Seg \mid \dots\}$ (def. of simple polygon)
 $Face = \{(c, H) \mid c \in Cycle, H \subset Cycle, \text{ such that } \dots\}$
 $D_{\text{region}} = \{F \subset Face \mid f_1, f_2 \in F \wedge f_1 \neq f_2$
 $\Rightarrow \text{edge-disjoint}(f_1, f_2)\}$

Moving Objects Databases – History of Movement

- Type System



Moving Objects Databases – History of Movement

- Type System as a ***signature***

<i>Type constructor</i>	<i>Signature</i>
<u>int</u> , <u>real</u> , <u>string</u> , <u>bool</u>	→ BASE
<u>point</u> , <u>points</u> , <u>line</u> , <u>region</u>	→ SPATIAL
<u>instant</u>	→ TIME
<u>range</u>	BASE \cup TIME → RANGE
<u>moving</u> , <u>intime</u>	BASE \cup SPATIAL → TEMPORAL

- A signature has ***sorts*** and ***operations***
- Sorts – collections of types
- Operations – type constructors

Moving Objects Databases – History of Movement

- Operations and QL – Design Goals
 - As generic as possible
 - Consistency between operations on non-temporal and temporal types
 - ***point x point*** ***→ real*** ***distance***
 - ***mpoint x mpoint*** ***→ mreal*** ***mdistance***
- 1. Define operations on non-temporal types
- 2. "Lift" them all to temporal types – lifting
- 3. Add specific operations for temporal types
- 4. To obtain a powerful query language, it is necessary to include operations from various domains:
 - *Simple set theory*
 - *FOL*
 - *Order relationships*
 - *Topology*
 - *Metric spaces*
 - ...

Moving Objects Databases – History of Movement

- Operations on Non-Temporal Types

Generic view: *point* and *point set* in some space

		π	σ
	Space	point type	point set type
	Integer	<u>int</u>	<u>range(int)</u>
	Real	<u>real</u>	<u>range(real)</u>
	Bool	<u>bool</u>	<u>range(bool)</u>
	String	<u>string</u>	<u>range(string)</u>
	Time	<u>instant</u>	<u>periods</u>
1D Spaces			
2D Space	2D	<u>point</u>	<u>points, line, region</u>

$\pi \times \sigma$	\rightarrow <u>bool</u>	inside	instantiates to
<u>int</u> \times <u>range(int)</u>	\rightarrow <u>bool</u>	inside	
<u>bool</u> \times <u>range(bool)</u>	\rightarrow <u>bool</u>		
<u>instant</u> \times <u>periods</u>	\rightarrow <u>bool</u>		
<u>point</u> \times <u>line</u>	\rightarrow <u>bool</u>		etc.

Moving Objects Databases – History of Movement

- Operations on Non-Temporal Types ...

Class	Operations
Predicates	isempty =, /=, intersects, inside <, <=, >=, >, before touches, attached, overlaps, on_border, in_interior
Set Operations	intersection, union, minus crossings, touch_points, common_border
Aggregation	min, max, avg, center, single
Numeric	no_components, size, perimeter, duration, length, area
Distance and Direction	distance, direction
Base Type Specific	and, or, not

Moving Objects Databases – History of Movement

Operations on Non-Temporal Types - Predicates

Unary

Operation	Signature	Semantics
isempty[undefined]	$\pi \rightarrow \underline{bool}$	$u = \perp$
	$\sigma \rightarrow \underline{bool}$	$U = \emptyset$

Binary



	Sets	Order (1D Spaces)	Topology
<i>point vs. point</i>	$u = v, u \neq v$	$u < v, u \leq v,$ $u \geq v, u > v$	
<i>point set vs. point set</i>	$U = V, U \neq V$ $U \cap V \neq \emptyset$ (intersects) $U \subseteq V$ (inside)	$\forall u \in U, \forall v \in V : u \leq v$ (before)	$\partial U \cap \partial V \neq \emptyset$ (touches) $\partial U \cap V^\circ \neq \emptyset$ (attached) $U^\circ \cap V^\circ \neq \emptyset$ (overlaps)
<i>point vs. point set</i>	$u \in V$ (inside)	$\forall u \in U : u \leq v$ (before) $\forall v \in V : u \leq v$	$u \in \partial V$ (on_border) $u \in V^\circ$ (in_interior)

Moving Objects Databases – History of Movement

- Operations on Temporal Types
 - Values of temporal types (i.e. $moving(\alpha)$) are partial functions

$$f: A_{\text{instant}} \rightarrow A_{\alpha}$$

- Projection to Domain and Range
 - deftime** - define set of time intervals when a temporal function is defined
 - rangevalues** - define set of time intervals when a temporal function is defined



Moving Objects Databases – History of Movement

- Operations on Temporal Types

<i>moving(α)</i>	→	<i>periods</i>	<i>deftime</i>
<i>moving(α)</i>	→	<i>range(α)</i>	<i>rangevalues</i>
<i>moving(point)</i>	→	<i>points</i>	<i>locations</i>
<i>moving(points)</i>	→	<i>points</i>	<i>locations</i>
<i>moving(point)</i>	→	<i>line</i>	<i>trajectory</i>
<i>moving(points)</i>	→	<i>line</i>	<i>trajectory</i>
<i>moving(line)</i>	→	<i>line</i>	<i>routes</i>
<i>moving(line)</i>	→	<i>region</i>	<i>traversed</i>
<i>moving(region)</i>	→	<i>region</i>	<i>traversed</i>

Moving Objects Databases – History of Movement

- Operations on Temporal Types

How large was the area of Austria affected by hurricane “Lizzy” ?

```
LET Austria = ...
```

```
LET Lizzy = ELEMENT (SELECT id  
FROM weather WHERE id = "Lizzy");  
area(intersection(traversed(Lizzy), Austria);
```

Where was flight "OS 7052" while hurricane “Lizzy” was over Austria ?

```
LET Austria = ...
```


```
trajectory(  
    atperiods      (ELEMENT(SELECT route FROM flight  
                           where id = "OS 7052"),  
    deftime(at(Lizzy, Austria)));
```

Moving Objects Databases – History of Movement


- Implementation

- Based on discrete model based on ***sliced representation***:

Temporal development of the value of type α by decomposing the time dimension into a set of disjoint time intervals (“slices”) such that within each slice the development can be described by some “simple” function.



moving(real)



moving(point)

References

- Allen J.F.: "*Maintaining Knowledge about Temporal Intervals*", Communications of the ACM,, 26(11), pp. 832-843. ACM Press, 1983.
- Güting R.H., Schneider M.: "*Moving Objects Databases*", Morgan Kaufmann Publishers, 2005.
- Koubarakis M., Frank A.U., Grumbach S, Güting R.H, Jensen C.S., Lorentzos N.A., Manolopoulos Y., Nardelli E., Pernici B., Schek H.J., Scholl M., Theodoulidis B., Tryfona N. : "*Spatio-Temporal Databases: The CHOROCHRONOS Approach*", LNCS 2520, Springer, 2003
- Snodgrass R. T.: "*Developing Time-Oriented Database Applications in SQL*", Morgan Kaufmann Publishers, 2000
- Steiner A.: "*A Generalisation Approach to Temporal Data Models and their Implementations*", Swiss Federal Institute of Technology, 1998
- Zaniolo C., Ceri S., Faloutsos C., Snodgrass R.T., Zicari, R.: "*Advanced Database Systems*", Morgan Kaufmann Publishers, 1997
- TSQL2: <http://www.cs.arizona.edu/~rts/tsql2.html>
- SQL3 <http://www.cs.arizona.edu/~rts/sql3.html>
- TimeCenter: <http://timecenter.cs.aau.dk/index.htm>